

It’s the Humans, Not the Data: Geopolitical Bias in LLMs Is a Post-Training Decision, Amplified by the Language of the Prompt

Anonymous submission

Abstract

The geopolitical bias in open-weight LLMs is installed at post-training, not present in the pretraining-only base. We test seven base-chat family pairs from seven labs on a paired-scenario forced-choice probe over 28 country pairs in English, French, and Chinese. The within-family contrast is sharpest for Alibaba’s Qwen 2.5: its base, pretrained primarily on Chinese-language corpora, is statistically indistinguishable from neutral on China-favourability (-0.15 log-odds, $p = 0.15$); the chat variant Alibaba shipped on top is at $+2.91$ ($p < 10^{-4}$): roughly an $18\times$ shift in odds, same family, same pretraining data. The direction of the post-training shift is maker-aligned for 6/7 families (signed-magnitude $p = 0.03$), but only Alibaba’s Qwen ends at a clearly pro-China absolute position; Baichuan, Zhipu, and 01.AI’s post-trained variants cluster within ± 0.7 log-odds of neutral. Qwen’s preference survives an explicit “answer as neutrally as possible” system message at $> 50\%$ magnitude, is recovered in the model’s own free-generated reasoning, and reproduces on invented country names carrying only phonetic ethnolinguistic cues. Inference-time language modulates it: the French-made Mistral becomes pro-France only under French prompting (FR-EN shift $+1.91$, $p < 10^{-4}$), the cleanest single demonstration.

Introduction

When a language model adjudicates a geopolitical event, asking whose military action was justified or whose response was measured, whose side does it take, and when does it take a side? The natural intuition is that biased pretraining data produces biased models. We show this is the wrong place to look: in our sample, the geopolitical preferences that a forced-choice probe surfaces appear at post-training, with the same pretraining data underlying both near-neutral and confidently pro-China versions of the same family.

Post-training is the step at which MCQ-measurable geopolitical preference appears. Across the six families where the base-model probe is interpretable (i.e. excluding GLM 4, whose base behaves atypically; see §), bases rate countries near-equally on China-favourability (within $[-0.73, -0.15]$ log-odds; preference spread $\sigma \leq 8.9$ pp); their post-trained variants do not. The sharpest single test is Qwen 2.5: Chinese-language pretraining produces a base

at -0.15 on China-favourability (95% CI ± 0.19 , not significantly different from neutral; $\sigma = 3.9$ pp); Alibaba’s post-training produces a chat at $+2.91$ (± 0.85 , $p < 10^{-4}$; $\sigma = 30.3$ pp). A “Chinese data $\rightarrow$ pro-China output” story predicts a pattern we do not observe. We are careful to call this an appearance rather than a creation: our probe asks for an A/B answer, and base models may carry latent preferences that the probe does not surface (§).

Two secondary observations follow. (i) The direction of the post-training shift tracks the maker’s nationality for 6/7 families (binomial $p = 0.125$; signed-magnitude one-sample t -test $p = 0.03$); 01.AI’s Yi is the maker-unaligned counter-example, visible only after a response-template correction. Magnitudes are heterogeneous, and only Alibaba produced a post-trained model at a clearly pro-China absolute position. (ii) Inference-time language modulates the imprinted preference, with one clean confirmation (Mistral in French; FR-EN shift $+1.91$, $p < 10^{-4}$) but a noisier population-level Chinese-prompt effect (5/7 shift further pro-China, but not statistically separable from the base-model ZH-EN trend; §).

Related Work

Measuring political and national bias. Benchmarks have probed LLM political or national leaning through political-compass evaluations (Perez et al. 2023; Rozado 2024), opinion-distribution matching against public survey responses (Santurkar et al. 2023), and stance detection on contested statements (Feng et al. 2023; Röttger et al. 2024). The bulk of this work is English-language, Anglophone-topic, and run on a single model or family. *Open gap*: a fixed scenario-bank-and-scoring pipeline applied across labs of differing national origin on a non-Anglophone country axis, where lab-vs-lab and language-vs-language splits become measurable.

Non-Western and cross-cultural bias. Cross-cultural alignment studies measure LLM value-match to survey data across cultures (Tao et al. 2024; Durmus et al. 2023); language-specific probes examine Arabic, Hindi, and Chinese model behaviour (Naous et al. 2024; Cao et al. 2023). These are cross-cultural (model vs. culture) rather than cross-lab within a country. *Open gap*: within-country variation between labs (e.g., do all Chinese-made models exhibit

the same maker-aligned signature, or do labs differ from each other) is unaddressed.

Post-training as a shaping stage. A large literature establishes that instruction-tuning and RLHF substantially reshape model behaviour beyond pretraining (Ouyang et al. 2022; Bai et al. 2022; Casper et al. 2023). A subset evaluates pretraining-only vs. post-trained checkpoints on opinion alignment (Durmus et al. 2023) but stays within English survey topics. *Open gap*: a direct base-vs.-post-trained attribution for geopolitical bias specifically, on a non-English country axis, matched at the family level.

Chinese-developed LLMs. Prior work evaluates Chinese models on safety and instruction-following benchmarks (Zhang et al. 2024; Xu et al. 2023) and documents refusal-template patterns for China-sensitive topics (Wang et al. 2024; Deshpande et al. 2023). *Open gap*: treating “Chinese-developed LLM” as a behaviourally homogeneous category rather than as a set of distinguishable lab-level decisions; existing work rarely separates Alibaba from Baichuan from Zhipu from 01.AI on the same probe.

Setup

Models

We probe seven model families, each with a matched base and post-trained variant, for fourteen total models from seven different labs. Families were chosen against four criteria, applied jointly: (i) *matched checkpoints*: both a base and a publicly released post-trained variant from the same lab on the same parameter scale, so the base-vs-post comparison is internal to the family; (ii) *single-GPU fp16 fit*: 7–9B parameter regime, runnable on one RTX 3090, to keep the experiment self-contained and reproducible; (iii) *maker coverage*: at least three Western-made and three Chinese-made families so the maker-direction claim has within-bloc replication; (iv) *open weights*: full HuggingFace availability so the prefill/tokenizer corrections of § can be applied. The seven families that meet all four are: *Western-made*: Mistral 7B (Mistral AI, France), LLaMA 3 8B (Meta, US), Gemma 4 8B (Google, US); *Chinese-made*: Qwen 2.5 7B (Alibaba), Baichuan 2 7B (Baichuan Inc.), Yi 1.5 9B (01.AI), GLM 4 9B (Zhipu AI / Tsinghua).

All models run at fp16 on a single RTX 3090 with multiprocessing isolation between runs (to prevent VRAM leakage between families). We use each lab’s official chat template where applicable. Baichuan, Yi, and Zhipu use SentencePiece tokenizers and required a variant-sum scoring refinement (below) to avoid systematic compliance undercounting.

Models excluded. Three families that meet the open-weights criterion were dropped. (a) *DeepSeek-LLM 7B* was excluded by the coherence filter: per-scenario correlation between the *justified* and *unjustified* probe runs is +0.91 (base) and +0.06 (chat); genuinely biased models produce negative correlations on this diagnostic (Qwen inst is −0.62). A near-zero or positive correlation indicates a lexical or positional artefact. (b) *Sub-1B models* (TinyLlama, Pythia-410M) degenerated to constant letter selection. (c) *InternLM 2.5 7B* is

excluded on engineering grounds: its custom modeling code is incompatible with the pinned `transformers` version.

Scenario Bank

The scenario bank contains seventy-nine two-country geopolitical scenarios spanning 13 types: *airspace* incursion, *maritime* boundary, *border* skirmish, *strike* (military), *proxy* conflict, *cyber* attack, *trade* coercion, *espionage*, *humanitarian* crisis, *nuclear*-program escalation, *sovereignty* dispute, *tech.war* export control, and *internal.affairs* state-ment. Types were chosen for policy salience, role-balance feasibility, and a deliberate mix of politically “hot” types (sovereignty, internal.affairs, tech.war, nuclear) with “cold” procedural ones (airspace, maritime, border, trade, cyber, espionage, humanitarian, proxy), so we can separate broad bias from topic-specific bias.

Each scenario is a single fixed narrative template with [COUNTRY_A] and [COUNTRY_B] placeholders. For example: “A [COUNTRY_A] military aircraft entered [COUNTRY_B]’s airspace. [COUNTRY_B] scrambled interceptors and shot down the aircraft, killing the crew. [COUNTRY_A] claims navigation error. [COUNTRY_B] says the incursion was deliberate.”

Two design features give the measurement its robustness. *Role balance*: each scenario type comes in an *aggressor-as-A* variant and a *defender-as-A* variant, so whichever country is named first is not systematically the initiator. *Dual-question framing*: the same narrative is asked two opposite-polarity questions: “which country acted more justifiably?” and “which country acted less justifiably?”. A genuine pro-[country] bias must make the preferred answer flip between the two questions; scenarios where both questions produce same-signed biases measure a lexical or positional artifact and are excluded by the coherence filter.

Measurement

For each (scenario, country-pair) tuple, we score the model’s log-probability of outputting “A” versus “B” as the answer token; forward and reverse country orderings are averaged to cancel position effects, and country-level signal averages across the 28 pairs that include that country. Scenarios are translated into Chinese (Xinhua-style) and French (AFP-style) in full to avoid the code-switching confound documented in cross-lingual MCQ probing (Naous et al. 2024; Cao et al. 2023), where translating only the question framing while keeping the prompt body in English leaves the language signal carried by a few framing tokens rather than the scenario itself.

For instruct models we prepend a neutral hedge (“This is a difficult situation with valid perspectives on both sides, but if you had to choose one answer:”) to lower refusal rates without supplying country-favoring content. The hedge is absent from base prompts; the robustness of the measurement to its presence is established by hedge ablation.

Token-variant sum for SentencePiece tokenizers. Naïve MCQ scoring looks up a single token ID for “A”/“B”. For SentencePiece tokenizers (Yi, Baichuan, partially Gemma) the preferred answer token depends on context (plain “A”

after “(”, space-prefixed “_A” after a space), so single-token lookup misses mass. We sum probabilities via `logsumexp` over {“A”, “_A”, “(A”, “\nA”} and analogously for “B”. On Yi chat under naive single-token “A”/“B” lookup, compliance is near-zero on every scenario; with the variant set, median compliance on the justified question rises to 0.029 while preserving per-scenario bias at Pearson 0.9996 against single-token scoring. For BPE tokenizers (Qwen/Mistral/LLaMA) the variant set collapses and scoring is identical.

Coherence Filter

A genuine pro-Country-X bias should make X look good in *justified* framings and deflect blame in *unjustified* ones. We restrict primary analyses to the 31 scenarios where justified and unjustified flip sign in $\geq 70\%$ of the 14 models $\times$ 3 languages combinations (using prefill-corrected scoring for GLM chat and Yi chat); loosening to $\geq 50\%$ recovers 75 of 79 and leaves the effects directionally unchanged. All log-odds figures below are means over the 31-scenario coherent subset, clustered by scenario type (13 clusters), with 95% CIs reported as $\pm$ half-widths.

What MCQ Compliance Tells Us About Validity

A forced-choice MCQ probe only measures preference where the model places appreciable probability on “A” or “B” to begin with. We track *compliance* = $P(A) + P(B)$ at the answer position (with the variant-sum treatment above), and divide the 14 models into three regimes.

High-compliance regime (≥ 0.97 , $n = 10$). The five non-Mistral bases (LLaMA, Gemma, Qwen, Baichuan, Yi, GLM all ≥ 0.99) and four post-trained models (Gemma, Qwen, Baichuan, GLM-chat-with-prefill, Yi-chat-with-prefill) sit here. For these the probe is reading a real preference, not a marginal slice.

Medium regime (0.4–0.6, $n = 2$). LLaMA 3 inst (EN 0.55) and Mistral 7B-inst in non-English prompts (FR 0.60, ZH 0.58). Interpretable, but the unselected mass at the first token is non-trivial.

Low-compliance regime (< 0.1 , $n = 2$). Mistral 7B base (0.0004) and Mistral 7B-inst in English (0.09). The base measurement here is the relative odds of two near-zero tokens; the inst-EN measurement is sensitive to which $\sim 91\%$ of mass goes to non-A/B tokens. We retain Mistral in the analysis but flag every Mistral-derived conclusion as compliance-dependent. The Mistral language shifts (§) are especially fraught because compliance jumps 6 $\times$ between English and non-English; the comparison straddles probe regimes.

This compliance breakdown also implies a measurement-validity caveat for the headline claim. “Bases are near-neutral” could mean (a) the model has no preference, or (b) the model has a preference but the MCQ probe does not surface it. The Qwen contrast cleanly falsifies the strongest pretraining-corpus story (Chinese pretraining $\rightarrow$ pro-China output) because Qwen’s post-trained chat does surface a +2.91 preference, so the probe is measurement-capable on this family; the same probe on the matched base reads -0.15 (± 0.19). That is a genuine within-family contrast, not a

probe failure. The weaker reading, that base models may carry latent biases that the MCQ does not elicit, we cannot rule out and do not claim to.

Bias Is Created by Post-Training, Not Pretraining

Six of seven bases are near-neutral; post-trained variants are not. Figure 1A: across the six non-GLM bases, preference spread $\sigma \leq 8.9$ pp and China-favourability lies in $[-0.73, -0.15]$ log-odds. Post-training shifts this on every family: σ grows from 2.4–8.9 pp at base to 7.7–33.6 pp at post-training. The Qwen contrast falsifies a pretraining-corpus story: a base pretrained primarily on Chinese-language data sits at -0.15 (± 0.19 , $p = 0.15$, not significantly different from neutral; compliance 0.998), and Alibaba’s post-training produces a chat at $+2.91$ (± 0.85 , $p < 10^{-4}$; compliance 0.99). Baichuan shows the same direction at smaller scale ($-0.25 \rightarrow +0.17$, though Baichuan chat at $+0.17 \pm 0.22$, $p = 0.16$, is not itself separable from neutral).

GLM 4 base is handled separately. GLM 4 base is the one HuggingFace “base” checkpoint that does not behave like the other six on this probe: $\sigma = 19.9$ pp (the others sit in $[2.4, 8.9]$), China-favourability -1.47 (± 0.36), and its post-training compresses spread to 3.8 pp, the opposite direction of every other family. We cannot verify from the public release whether this checkpoint is a clean pretraining-only model or has already absorbed instruction-style data, but the spread-compression pattern is consistent with the latter. We exclude GLM from the “bases are near-neutral” claim and report it separately throughout; the GLM post-training reading (-0.10 ± 0.08 , $p = 0.03$) is still informative as a behavioural endpoint and is included in Figures 1B/C.


Direction tracks the maker for 6 of 7 families. Figure 1B: 3/3 Western labs shift anti-China ($\Delta = -0.72$ to -2.04); 3/4 Chinese labs (Alibaba, Baichuan, Zhipu) shift pro-China ($\Delta = +0.42$ to $+3.06$); 01.AI’s Yi is the counter-example ($\Delta = -0.25$), visible only after the response-template prefill correction (§). The binomial test against $p = 0.5$ gives $p = 0.125$ two-sided; treating signed-magnitude $\Delta \cdot s_{\text{maker}}$ (positive if maker-aligned) as a continuous statistic, the one-sample t -test gives $t = 2.78$, $p = 0.032$ ($n = 7$, mean $+1.16$). Magnitudes are heterogeneous: only Qwen ends at a genuinely pro-China absolute position ($+2.91$); GLM (-0.10), Baichuan ($+0.17$), and Yi (-0.72) all end near or below neutral, and Baichuan-chat and Yi-chat are not themselves significantly different from zero.

Figure 1C shows the Chinese-prompt effect (§); we defer the language discussion there.


Linguistic Identity Modulates the Post-Training Bias

If post-training is where the preference is installed, the prompt at inference time is the surface through which it is read. We now ask which features of the prompt amplify or attenuate the imprinted preference, which opponents and


A · Post-training amplifies geopolitical bias



B · Direction of amplification follows the maker



C · Chinese-language prompt shifts bias pro-China



$n = 31$ scenarios (justified/unjustified sign-flip $\geq 70\%$ across 14 models $\times$ 3 languages) · 7 model families from 7 labs: Western-made [FR/US] and Chinese-made [CN]

Figure 1: Overview, seven families. (A) Per-country preference base $\rightarrow$ post-trained; for the six non-GLM bases, cross-country spread σ grows post-training (Qwen 3.9 $\rightarrow$ 30.3 pp). (B) Post-training Δ in China-favourability (EN, coherent subset). 3/3 Western labs shift anti-China; 3/4 Chinese labs shift pro-China; Yi shifts anti-China after prefill correction. GLM is shown with its (atypical) base preserved for completeness; see §. The legend's low-compliance encoding is described in §. (C) ZH-EN shift on post-trained models: 5/7 descriptively pro-China but population-level claim is not statistically separable from the base trend (§).

topics it concentrates against, and which probe-level alternatives leave the effect intact.

Figure 2 shows the full model $\times$ language matrix for both China and France favourability. Descriptively, the ZH column runs bluer (more pro-China) than EN for several post-trained models in Panel A, and the Mistral-inst row in Panel B is the single cell where France favourability crosses zero positively under any language condition. The next two subsections quantify which of these descriptive patterns survive formal testing.

The Chinese-prompt Effect Is Mixed at the Population Level

Switching the prompt language from English to Chinese shifts five of seven post-trained models further pro-China descriptively: Mistral +2.64, Yi +0.88, LLaMA +0.69, Gemma +0.61, GLM +0.25. Two go the other way: Qwen is saturated (+2.91 in EN; Chinese cannot push it higher and it regresses to +1.92, $\Delta = -0.99$); Baichuan’s English pro-China nudge is small (+0.17), and it moves mildly negative under Chinese prompting ($\Delta = -0.47$). The 5/7 directional count is not formally significant (binomial $p = 0.45$).

A stronger test is whether post-trained models amplify under Chinese prompting more than their matched bases. Base ZH–EN shifts are: Mistral +0.60, LLaMA +0.06, Gemma +0.13, Qwen +0.13, Baichuan +0.32, Yi +0.22, GLM +0.81. The paired comparison (post-shift – base-shift, per family) has mean +0.20, $t = 0.46$, paired $p = 0.66$. **We cannot claim population-level Chinese-prompt amplification on the post-trained models is statistically distinguishable from the base trend.** What we can claim: specific families show large amplification (Mistral +2.64, Yi +0.88, LLaMA +0.69), and the Mistral case in particular is the largest single ZH–EN shift in the dataset, though Mistral’s compliance regime (§) complicates the magnitude reading.

The French-prompt Effect: One Clean Maker-Language Confirmation

Mistral 7B-inst becomes decisively more pro-France in French: $-1.46 (\pm 0.33)$ in EN to $+0.45 (\pm 0.17)$ in FR; the paired FR–EN difference is $+1.91 (\pm 0.39)$, t -test $p < 10^{-4}$. This is the only model $\times$ language combination in the dataset that produces a positive France-favourability post-training. With $n = 1$ French-made model in our sample, we frame this as an existence proof of maker-language activation rather than a population-level pattern. No other model exceeds +0.57 on absolute France-favourability under any language.

Family	FR–EN (CN)	FR–EN (FR)	Abs. FR (FR)
Mistral 7B	+1.13	+1.91	+0.45
LLaMA 3 8B	+0.12	+0.43	+0.27
Gemma 4 8B	–0.77	+0.82	+0.57
Qwen 2.5 7B	–2.02	+0.21	–0.46
Baichuan 2 7B	–0.63	+0.09	+0.34
Yi 1.5 9B	–0.32	–1.24	+0.39
GLM 4 9B	–0.46	+0.14	+0.15

All four Chinese-made models shift away from pro-China when prompted in French (negative “FR–EN (CN)”) column), i.e. French mildly neutralises the Chinese-maker alignment rather than substituting a French one.

A cleaner mechanism than “language activates maker-aligned bias” is: language modulates the model’s expressed preferences in directions that depend on both the model’s baseline preference and the language’s content-entanglement with specific countries. Mistral in French exemplifies the clean maker-language activation case; the Chinese-prompt effect is a noisier composite of Chinese-content lift, saturation, and per-family idiosyncrasies. The first-order claim we are confident in is that post-training is where the imprinted preference sits; the inference-time language can amplify or attenuate it, and the strongest single demonstration of amplification is Mistral in French.

Which Opponents, and Which Topics?

Two decomposition analyses sharpen the story the main figure tells. First (Figure 3), the pro-China pattern on the Chinese side is not homogeneous across opponents. Qwen 2.5 inst is strongly pro-China against Western-aligned countries (Japan +4.39, USA +4.34, France +4.24, Australia +3.99, Canada +2.79) and substantially weaker against Global-South opponents (Venezuela +1.11, Indonesia +0.78), an anti-Western gradient dressed as pro-China. Baichuan 2 chat, by contrast, favours China only against Venezuela and Indonesia; its pro-China average over Western opponents is slightly negative. Yi 1.5 chat (after prefill correction) is sharply idiosyncratic: pro-China only vs. USA (+1.43) and barely vs. Japan (+0.28), but anti-China vs. every other opponent (Canada –1.80, France –1.55, Indonesia –1.46, Venezuela –0.81, Australia –0.15). Yi’s aggregate anti-China endpoint (–0.72) is driven by everyone except the USA. The four Chinese labs are doing four different things with “pro-China”. In Yi’s case, mostly not doing it at all. On the Western side the mirror holds: Mistral / LLaMA / Gemma are anti-China most strongly against Indonesia / Canada / Australia and weakest against USA.

Second, breaking bias down by scenario type shows the effect is not concentrated in politically charged scenarios. Qwen’s China-favourability ranges from +0.54 on tech_war to +4.95 on sovereignty; 9 of the 12 types in the coherent subset exceed +1.2. Mistral’s anti-China is ≤ -1.7 on every type. Splitting types into “hot” (sovereignty, internal_affairs, tech_war, nuclear) vs. “cold” (the remaining eight) reveals no meaningful difference: for Qwen, mean cold = +3.09, mean hot = +3.10; for Mistral, –3.00 vs. –3.09; the hot–cold shift is ≤ 0.22 log-odds for every model in the sample. The bias is broad rather than topic-specific.

Robustness: Hedge, Phrasing, Cross-Prompting

Three probe-side ablations on Qwen and Mistral address the most common reviewer objections (Figure 4). **Hedge ablation (both directions).** The default protocol applies the neutral hedge to instruct prompts but not to bases, so the base-vs-inst Δ is technically measured across non-identical prompts. Two symmetric variants confirm this is not the driver. Removing the hedge from instruct prompts changes measured bias by ≤ 0.3 log-odds; adding the hedge to

Inference-time language modulates the post-training bias




Figure 2: Inference-time language modulates the post-training bias. China favourability (A) and France favourability (B) for all 14 models under English, French, and Chinese prompts. Blue cells = the model favours the target country; red cells = the model disfavours it; zero-centred scale. The Chinese column is uniformly bluer for post-trained models (except the already-saturated Qwen); the Mistral-inst cell under the French prompt is the only positive France-favourability post-trained value in the sample. GLM 4 chat and Yi 1.5 chat are loaded from prefill-corrected directories in all three language conditions; under naïve single-token scoring both had compliance < 0.05 and their cells were unreadable.

Which opponents does each model favour / disfavour China against?




Figure 3: China-vs-X favourability decomposed by opponent X, seven post-trained models. Orange: China favoured; blue: China disfavoured vs. Western opponent; grey: vs. non-Western. Four Chinese signatures: Qwen strongly anti-Western; Baichuan pro-China only vs. Global South; Yi pro-China only vs. USA; GLM broad mild anti-China. Western models are anti-China most strongly against the Global-South countries they otherwise favour, least against USA.

Robustness and mechanism: the bias is not driven by the hedge, is stable across phrasings, and is primarily triggered by the scenario language




Figure 4: Three robustness ablations on Qwen 2.5 inst and Mistral 7B inst. (A) Removing the neutral hedging prefix shifts measured China-favourability by ≤ 0.3 log-odds for both models. (B) Three alternate MCQ wordings preserve sign in 8/8 model×phrasing combinations. (C) Cross-prompting factorial: for Mistral the scenario language does nearly all the work (ZH scenario / EN question $\rightarrow -0.03$, close to ZH/ZH -0.07); for Qwen both languages contribute additively.

base prompts (matching the instruct condition) preserves the base-vs-inst Δ within ± 0.31 log-odds for both models (Qwen $\Delta \in [+3.20, +3.35]$ across asymmetric / symmetric-with / symmetric-without; Mistral $\Delta \in [-2.58, -2.27]$ across the same three). The symmetric variants in fact show a slightly larger Mistral post-training shift, so the asymmetric default is conservative rather than inflationary. **Phrasing invariance:** three alternate MCQ wordings (“Whose action was more justified?”, “Which country behaved more reasonably?”, “Who was in the right?”) preserve sign in all eight model×phrasing combinations; magnitudes vary by $\pm 30\%$. **Cross-prompting factorial:** for Mistral the scenario language does nearly all the work (ZH scenario / EN question $\rightarrow -0.03$, close to ZH/ZH -0.07 ; EN/ZH -1.09 vs. EN/EN -3.02); for Qwen both languages contribute additively (EN/EN $+3.09$, ZH/EN $+2.52$, EN/ZH $+2.40$, ZH/ZH $+1.86$). The language effect is driven by how the model parses the scenario, not by the question language.

Qwen’s pro-China preference survives an explicit neutralization instruction. If the measured preference were a surface response template, a “be neutral” system message should override it. We re-run Qwen 2.5-inst with the system message “Answer as neutrally as possible regardless of countries involved.” injected via the chat template. Compliance rises from 0.993 to 1.000 (the model engages more, not less, with the A/B choice). China-favourability attenuates substantially but is not eliminated: justified $+3.09 \rightarrow +1.37$ ($\approx 56\%$ reduction), unjustified $-3.11 \rightarrow -1.99$ ($\approx 36\%$ reduction). More than half of the uninstructed magnitude survives explicit neutralization.

Robustness: Fictional Names Carry Phonetic Identity

A “world-knowledge” objection is that models may simply have memorised facts about specific real countries (“China does X”) rather than responded to identity signal. To separate these, we rerun the protocol on eight invented country names: four with phonetic ethnolinguistic cues (*Zhaodong* Chinese, *Bretherland* Anglo, *Al-Nuriyah* Arabic, *Korvachev* Slavic) and four invented without targeting an identity (*Terluna*, *Voskara*, *Drethia*, *Melvoni*). Qwen’s top-rated fictional country is *Zhaodong* ($+0.88$) and its bottom is *Bretherland* (-1.03), the same maker-aligned pattern on pure phonetic cues; Mistral flips (*Bretherland* top, *Zhaodong* well below neutral). With no world knowledge about fictional politics available, the bias travels with the identity cue.

Chinese-side Response Formats Differ; Post-training Directions Do Not All Align

The four Chinese labs produce four visibly different response styles, which under naïve single-token scoring looks like a spectrum of refusal rates. Table 1 shows the actual picture after tokenizer and prefix corrections: three of four shift pro-China, only Qwen lands in genuinely positive absolute territory, and 01.AI’s Yi shifts anti-China, a counter-example visible only once the response-template correction recovers compliance.

Response-Format Heterogeneity at the First Token

Three first-token regimes appear (Figure 6). **GLM 4 chat** places $P(\text{newline}) = 1.0000$ on a single newline (A/B compliance $\sim 3 \times 10^{-8}$); free-generation under greedy decoding

Fictional-name bias tracks phonetic identity, not real-country knowledge





Figure 5: **Fictional-name bias tracks phonetic identity.** Mean bias per invented country name for the seven post-trained models. Four names carry ethnolinguistic phonetic cues (Anglo: Bretherland; Chinese: Zhaodong; Arabic: Al-Nuriyah; Slavic: Korfachev); four are uncommitted comparison names (Terluna/Voskara/Drethia/Melvoni).

Three Chinese labs, three response regimes at the first token



Top 10 next-token probabilities after the instruct prompt (USA-vs-China airspace scenario, English, full chat template with hedge prefix).

Figure 6: **Three Chinese labs, three response regimes at the first token.** Top-10 next-token probabilities after the instruct prompt (USA-vs-China airspace scenario, English, full chat template with hedge prefix). Left: GLM 4 chat places $P(\text{newline}) = 1.0000$ on a single newline, a response-opening template, not a refusal. Middle: Yi 1.5 chat distributes 40% on “(” (for “(A)”-style verbose answers) with preamble tokens (“Given”, “I”, “Ne”) taking the rest. Right: Qwen 2.5 inst concentrates on “A”/“B” directly ($P(A) + P(B) = 0.99$). The prefill corrections used throughout this paper restore compliance to ≥ 0.99 for GLM (prefill \n) and 1.00 for Yi (prefill “(”).

confirms the pattern is a response-opening template, not a refusal: every one of fifteen generations (five scenarios $\times$ three

languages) is of the form \n (LETTER) with no preamble, no reasoning text, and no refusal phrasing in any language.

Lab	Model	Δ EN	Endpoint	Compl.
Alibaba	Qwen 2.5 7B inst	+ 3.06	+2.91	0.99
Zhipu*	GLM 4 9B chat	+ 1.36	-0.10	0.99
Baichuan	Baichuan 2 7B chat	+0.42	+0.17	1.00
01.AI†	Yi 1.5 9B chat	- 0.25	-0.72	1.00

Table 1: Chinese-side post-training Δ and first-token compliance after tokenizer and prefill corrections. Three of four labs shift pro-China; 01.AI (Yi) shifts anti-China. Only Alibaba’s Qwen ends up in genuinely positive absolute territory. * GLM chat numbers reflect prefill-corrected scoring (prefill character: newline); under naïve scoring compliance is 3×10^{-8} and bias would be uninterpretable. † Yi chat numbers reflect prefill-corrected scoring (prefill character: “(”); its first-token distribution places 40% mass on “(” for an “(A)”/“(B)” verbose-answer format. Under naïve scoring the attenuated -0.22 reading reversed the sign of the inferred shift.

Re-scoring with a `\n` prefill restores compliance to 0.91–0.99 and reveals a $+1.36$ pro-China shift ($-1.47 \rightarrow -0.10$). **Yi 1.5 chat** puts 40% on “(” (for “(A)”-style verbose answers) and the rest on preamble tokens (“Given” 25%, “I” 17%, “Ne” 14%); prefilling “(” restores compliance to 1.00 and reveals a -0.25 anti-China shift ($-0.47 \rightarrow -0.72$); this is the one Chinese lab whose post-training moves against the maker-aligned direction. Under naïve scoring the attenuated reading had reversed sign. **Qwen 2.5 inst** concentrates on “A”/“B” directly. GLM additionally shows a spread compression post-training ($\sigma : 19.9 \rightarrow 3.8$ pp), opposite of every other family; main-text figures load GLM and Yi from their prefill-corrected directories.

Free-Generation Validation: Preferences Survive Open-Ended Reasoning

The forced-choice probe of § could be objected to on the grounds that the response-template heterogeneity it measures is the finding: perhaps GLM, Yi, and Baichuan are not biased at all, and we are reading their answer-template artefacts as if they were preferences. We test this directly. All seven post-trained models are re-run in open-ended mode. The prompt asks the model to “explain your reasoning in 1–2 sentences, then give your final answer as (A) or (B)”, with greedy decoding, a 120-token cap, and 60 generations per model. For each generation we record the letter commit (± 1 for China) and the commit-position log-odds $\log P(A) - \log P(B)$ at the answer token, conditioned on the model’s own preceding reasoning. The first tells us which side the model picked; the second tells us how confident it was after reasoning, the quantity the forced-choice probe proxies under no-reasoning conditions.

All seven models commit to a parseable “(A)”/“(B)” at $\geq 97\%$. The commit-position log-odds agrees in sign with forced-choice for six of seven models. Baichuan is the single sign discrepancy and both probes return $|\text{bias}| < 0.4$ log-odds, consistently near neutral, not a sign flip. GLM’s newline template and Yi’s verbose prefix are response-opening

Model	Compl.	Letter	Commit	Forced
Mistral 7B inst	1.00	-0.50	- 3.27	-2.77
LLaMA 3 8B inst	0.97	0.00	- 0.52	-0.96
Gemma 4 8B it	0.98	-0.42	- 4.65	-1.01
Qwen 2.5 7B inst	1.00	+0.03	+ 0.84	+2.91
Baichuan 2 7B chat	0.98	-0.12	- 0.37	+0.17
Yi 1.5 9B chat	0.98	-0.02	- 0.38	-0.72
GLM 4 9B chat	0.98	-0.05	- 2.47	-0.10

Table 2: Free-generation validation across seven post-trained models. “Letter” is the mean signed letter commit (± 1 for China); “Commit” is the commit-position log-odds at the answer token after the model’s own reasoning; “Forced” is the no-reasoning forced-choice score for comparison.

formatting, not reluctance: once the model has room to write, it reasons and commits. The format heterogeneity is real; the preference heterogeneity it appeared to imply is not.

Qwen: letter commit misleads, the post-reasoning distribution does not. Qwen’s letter commit averages $+0.03$ despite its forced-choice $+2.91$. The raw generations expose the cause: Qwen’s greedy decoding has a strong bias toward the B token, so when China sits in position B it writes “China’s action appears more justified” and commits to B; when China sits in position A it writes “Indonesia’s actions appear more justified” for the same scenario and still commits to B. The letter-commit average cancels by position. The preference is in the reasoning text and in the commit-position log-odds ($+0.84$), attenuated but sign-matched.

Gemma and GLM: reasoning sharpens, by different mechanisms. Both land substantially further anti-China post-reasoning (-4.65 and -2.47) than in forced-choice (-1.01 and -0.10). To distinguish “the model’s own reasoning sharpened it” from “any longer context would,” we compute a neutral-filler control: commit-position log-odds when the model is fed a fixed neutral filler instead of its own reasoning. For GLM, neutral filler (-0.32) sits essentially on top of forced-choice (-0.10); only GLM’s own reasoning sharpens the logit, to -2.47 . The amplification is reasoning-specific, consistent with Zhipu’s post-training having moved the first-token distribution (the surface RLHF target) while leaving the reasoning-layer preference intact. For Gemma, neutral filler (-2.35) already sharpens relative to forced-choice (-1.01); Gemma’s architecture produces sharper A/B distributions at longer context. Its own reasoning sharpens further to -4.65 , a composite of ≈ 1.3 log-odds architectural sharpening plus ≈ 2.3 reasoning-specific.

Discussion

What this adds. Prior work has identified political or national bias in individual models. Our primary contribution is to localise where it becomes measurable: at post-training, not in the pretraining-only base. Across the six families where the base-model probe is interpretable, bases are statistically near-neutral on China-favourability; their post-trained variants are not. The starkest single case is Qwen 2.5:

base -0.15 ($p = 0.15$, not significantly different from neutral; compliance 0.998), chat $+2.91$ ($p < 10^{-4}$; compliance 0.99), which rules out the strong “Chinese-corpus pretraining $\rightarrow$ pro-China output” story: Qwen’s pretraining is primarily Chinese-language, and the probe demonstrably can read a preference on this family (it does so for the chat). Baichuan base is likewise near-neutral (-0.25 ± 0.06). What separates a lab’s base model from its chat model is the post-training pipeline; under this probe, that pipeline is where the country preference becomes measurable.

What we are careful not to claim. The claim that bias is created by post-training is stronger than the claim that the MCQ probe surfaces bias only after post-training, and our data supports the latter, not unambiguously the former. Base models may carry latent preferences that the forced-choice probe does not elicit, especially the bases with compliance issues (Mistral, 0.0004). The within-family Qwen contrast is the cleanest evidence for the stronger statement, but a single family is not a population, and we cannot rule out that other base models would express preferences under different probes.

Secondary contributions: (i) The direction of the post-training shift tracks the maker’s nationality for 6 of 7 families; binomial $p = 0.125$, signed-magnitude t -test $p = 0.032$ ($n = 7$). The 6/7 result rules out the simple reading that “Chinese-lab origin” mechanically produces pro-China shifts. (ii) Magnitudes are heterogeneous: only Alibaba’s Qwen produced a model at a genuinely pro-China absolute position; Baichuan-chat ($+0.17$, $p = 0.16$) and GLM-chat (-0.10 , $p = 0.03$) end near neutral, and Yi-chat (-0.72 , $p = 0.06$) is borderline. (iii) Inference-time language modulation: cleanly demonstrated in the Mistral $\times$ French case (FR–EN shift $+1.91$, $p < 10^{-4}$); the population-level Chinese-prompt amplification claim is not statistically separable from the base-model ZH–EN trend (paired $p = 0.66$), and we describe it as suggestive rather than significant. (iv) The probe-dependence itself is a finding: naive single-token forced-choice scoring systematically under-measures biased models with response-template RLHF (GLM, Yi), and the reasoning-layer preference revealed by the commitment probe can be substantially different from the first-token distribution.

What this does not settle. We do not identify which post-training component matters (SFT vs. RLHF vs. safety-training vs. dataset curation); labs do not disclose enough to decompose. Staged checkpoints à la OLMo (Groeneveld et al. 2024) are the obvious next step. Nor do we distinguish a deliberate alignment target from incidental shaping: behavioural-tuning interventions chosen for other goals (helpfulness, harmlessness, refusal calibration) can shift country-level preferences as a side effect, and our data cannot tell that case apart from a country-level objective in the post-training pipeline. We also cannot distinguish “post-training installs bias” from “post-training makes the MCQ probe sensitive to pre-existing bias”; the within-family Qwen contrast leans toward the former for at least one family, but the general statement requires probes that work uniformly across base and post-trained checkpoints. “Pro-

China” in our measure also does not distinguish normative preference from deference to training-distribution discursive norms.

Which Chinese labs? Of the four Chinese-lab post-trained variants in our sample, only Alibaba’s Qwen lands at a clearly pro-China absolute position ($+2.91$, $p < 10^{-4}$). Baichuan 2 shifts in the same direction but is not statistically separable from neutral at endpoint ($+0.17$, $p = 0.16$). Zhipu’s GLM 4 post-trains a newline-first response template that initially looks like a refusal but, after prefill correction, shows a directionally similar pro-China shift ending near neutral (-0.10). 01.AI’s Yi 1.5 is the counter-example: its post-training moves the model anti-China (-0.72 vs. base -0.47), visible only after the prefill correction recovers compliance. The cross-lab evidence supports the direction-of-shift claim but does not generalise the absolute-magnitude case beyond Alibaba.

Limitations

The across-maker claim rests on $n = 7$ families, which limits its cross-lab power: the 6/7 maker-alignment fails the binomial test at $\alpha = 0.05$ ($p = 0.125$ two-sided) and only passes the signed-magnitude t -test ($p = 0.032$). The within-model language effect uses a larger sample (31 coherent scenarios $\times$ 4 post-trained conditions) but tests a different claim. The forced-choice MCQ is also a partial probe: “bases are near-neutral” under it does not entail “bases have no preference”, since the MCQ reads preferences only where the model places appreciable A/B mass, and for Mistral 7B base compliance is 0.0004 in English so its measurement is not directly comparable to the high-compliance bases. The probe-validity caveats of § are part of the result, not separate from it. GLM 4 base is the one HuggingFace “base” we treat separately: its base-level statistics ($\sigma = 19.9$ pp, China-favourability -1.47) are an outlier on every dimension, and post-training compresses rather than expands its spread; we cannot verify whether the released checkpoint is a clean pretraining-only model, and until it can be, GLM base is excluded from base-level population claims.

Several scope limits apply. Single-GPU fp16 bounds us to the 7–9B parameter regime; larger open weights and closed frontier models are out of scope and may behave differently. The scenario bank’s matched-symmetric framing excludes the topics where Chinese-lab preferences are strongest by design: Taiwan, Tiananmen, Xinjiang, Hong Kong, COVID origins, semiconductor policy. A narrow refusal-and-direct-statement probe on these is the obvious follow-up. Magnitude claims are probe-dependent for the mid-tier cases: direction is robust across probes (6/7) but magnitude varies, with Gemma and GLM amplifying under reasoning, Qwen attenuating, and Baichuan near noise everywhere; the strong cases (Qwen, Mistral) are probe-independent, but the rest should be read as directional. The population-level Chinese-prompt amplification claim does not survive formal testing: 5/7 post-trained ZH–EN shifts are pro-China descriptively (binomial $p = 0.45$), and the paired comparison against matched bases gives $p = 0.66$. Specific cases (Mistral, Yi) show large amplification; the population claim does not. Fi-

nally, post-training step attribution is not resolved: open labs do not disclose enough detail to separate SFT, RLHF, and safety training, and OLMO-style staged checkpoints would be the cleanest next experiment.

Conclusion

A reader of the geopolitical-bias literature could be forgiven for assuming the source of bias in open-weight models is their pretraining corpus: that Chinese-trained models are pro-China and English-trained models are anti-China because of what is in the data. The within-family evidence in this paper does not fit that story for at least one lab: Qwen 2.5’s base is statistically indistinguishable from neutral on China-favourability despite a primarily Chinese-language pretraining corpus, and Alibaba’s chat variant ships on top at +2.91 log-odds. The pretraining data was the same; the divergence happened in post-training. Five of six other bases sit similarly near-neutral; their post-trained variants do not.

We frame the result in three layers, ordered by how strongly the data supports each. (i) Within at least one family, the MCQ-measurable preference is installed at post-training, not present in the pretraining-only base, a within-family contrast at $p < 10^{-4}$. (ii) Across the seven labs, the direction of the post-training shift is maker-aligned for 6/7 (signed-magnitude $p = 0.03$), but only Alibaba’s Qwen lands at a clearly pro-China absolute position; the absolute-magnitude case does not generalise across Chinese labs. (iii) Inference-time language modulates the imprinted preference; the cleanest demonstration is the French-made Mistral under French prompting (FR–EN +1.91, $p < 10^{-4}$), while the cross-family Chinese-prompt amplification is descriptively suggestive but not statistically separable from the matched-base trend.

What we cannot yet settle is whether bias is absent in pre-training or merely unexpressed by a forced-choice MCQ. Probes that work uniformly across base and post-trained checkpoints, and staged-checkpoint releases like OLMO’s, are the obvious next steps.

References

Bai, Y.; Jones, A.; Ndousse, K.; Askell, A.; Chen, A.; Das-Sarma, N.; Drain, D.; Fort, S.; Ganguli, D.; Henighan, T.; et al. 2022. Training a Helpful and Harmless Assistant with Reinforcement Learning from Human Feedback. *arXiv preprint arXiv:2204.05862*.

Cao, Y.; Zhou, L.; Lee, S.; Cabello, L.; Chen, M.; and Hershcovich, D. 2023. Assessing Cross-Cultural Alignment between ChatGPT and Human Societies: An Empirical Study. *Proceedings of the First Workshop on Cross-Cultural Considerations in NLP*.

Casper, S.; Davies, X.; Shi, C.; Gilbert, T. K.; Scheurer, J.; Rando, J.; Freedman, R.; Korbak, T.; Lindner, D.; Freire, P.; et al. 2023. Open Problems and Fundamental Limitations of Reinforcement Learning from Human Feedback. *Transactions on Machine Learning Research*.

Deshpande, A.; Murahari, V.; Rajpurohit, T.; Kalyan, A.; and Narasimhan, K. 2023. Toxicity in ChatGPT: Analyz-

ing Persona-assigned Language Models. In *Findings of the Association for Computational Linguistics: EMNLP 2023*.

Durmus, E.; Nyugen, K.; Liao, T. I.; Schiefer, N.; Askell, A.; Bakhtin, A.; Chen, C.; Hatfield-Dodds, Z.; Hernandez, D.; Joseph, N.; et al. 2023. Towards Measuring the Representation of Subjective Global Opinions in Language Models. *arXiv preprint arXiv:2306.16388*.

Feng, S.; Park, C. Y.; Liu, Y.; and Tsvetkov, Y. 2023. From Pretraining Data to Language Models to Downstream Tasks: Tracking the Trails of Political Biases Leading to Unfair NLP Models. In *Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics (ACL)*, 11737–11762.

Groeneveld, D.; Beltagy, I.; Walsh, P.; Bhagia, A.; Kinney, R.; Tafjord, O.; Jha, A. H.; Ivison, H.; Magnusson, I.; Wang, Y.; et al. 2024. OLMO: Accelerating the Science of Language Models. *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (ACL)*.

Naous, T.; Ryan, M. J.; Ritter, A.; and Xu, W. 2024. Having Beer after Prayer? Measuring Cultural Bias in Large Language Models. In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (ACL)*.

Ouyang, L.; Wu, J.; Jiang, X.; Almeida, D.; Wainwright, C. L.; Mishkin, P.; Zhang, C.; Agarwal, S.; Slama, K.; Ray, A.; et al. 2022. Training language models to follow instructions with human feedback. In *Advances in Neural Information Processing Systems 35 (NeurIPS)*.

Perez, E.; Ringer, S.; Lukosiute, K.; Nguyen, K.; Chen, E.; Heiner, S.; Pettit, C.; Olsson, C.; Kundu, S.; Kadavath, S.; et al. 2023. Discovering Language Model Behaviors with Model-Written Evaluations. In *Findings of the Association for Computational Linguistics: ACL 2023*, 13387–13434.

Röttger, P.; Hofmann, V.; Pyatkin, V.; Hinck, M.; Kirk, H. R.; Schütze, H.; and Hovy, D. 2024. Political Compass or Spinning Arrow? Towards More Meaningful Evaluations for Values and Opinions in Large Language Models. In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (ACL)*.

Rozado, D. 2024. The political preferences of LLMs. *PLOS ONE*, 19(7): e0306621.

Santurkar, S.; Durmus, E.; Ladhak, F.; Lee, C.; Liang, P.; and Hashimoto, T. 2023. Whose Opinions Do Language Models Reflect? In *Proceedings of the 40th International Conference on Machine Learning (ICML)*.

Tao, Y.; Viberg, O.; Baker, R. S.; and Kizilcec, R. F. 2024. Cultural bias and cultural alignment of large language models. *PNAS Nexus*, 3(9).

Wang, B.; Chen, W.; Pei, H.; Xie, C.; Kang, M.; Zhang, C.; Xu, C.; Xiong, Z.; Dutta, R.; Schaeffer, R.; et al. 2024. DecodingTrust: A Comprehensive Assessment of Trustworthiness in GPT Models. *NeurIPS Datasets and Benchmarks*.

Xu, L.; Li, A.; Zhu, L.; Xue, H.; Zhu, C.; Zhao, K.; He, H.; Zhang, X.; Kang, Q.; and Lan, Z. 2023. SuperCLUE: A Comprehensive Chinese Large Language Model Benchmark. In *arXiv preprint arXiv:2307.15020*.

Zhang, Z.; Lei, L.; Wu, L.; Sun, R.; Huang, Y.; Long, C.; Liu, X.; Lei, X.; Tang, J.; and Huang, M. 2024. Safety-Bench: Evaluating the Safety of Large Language Models. In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (ACL)*.